



SEOUL NATIONAL UNIVERSITY
Graduate School of Data Science

M3239.003100: Data Analysis and Visualization

Lecture 4

Univariate Analysis

Hyunwoo Park

Graduate School of Data Science

Seoul National University

Agenda

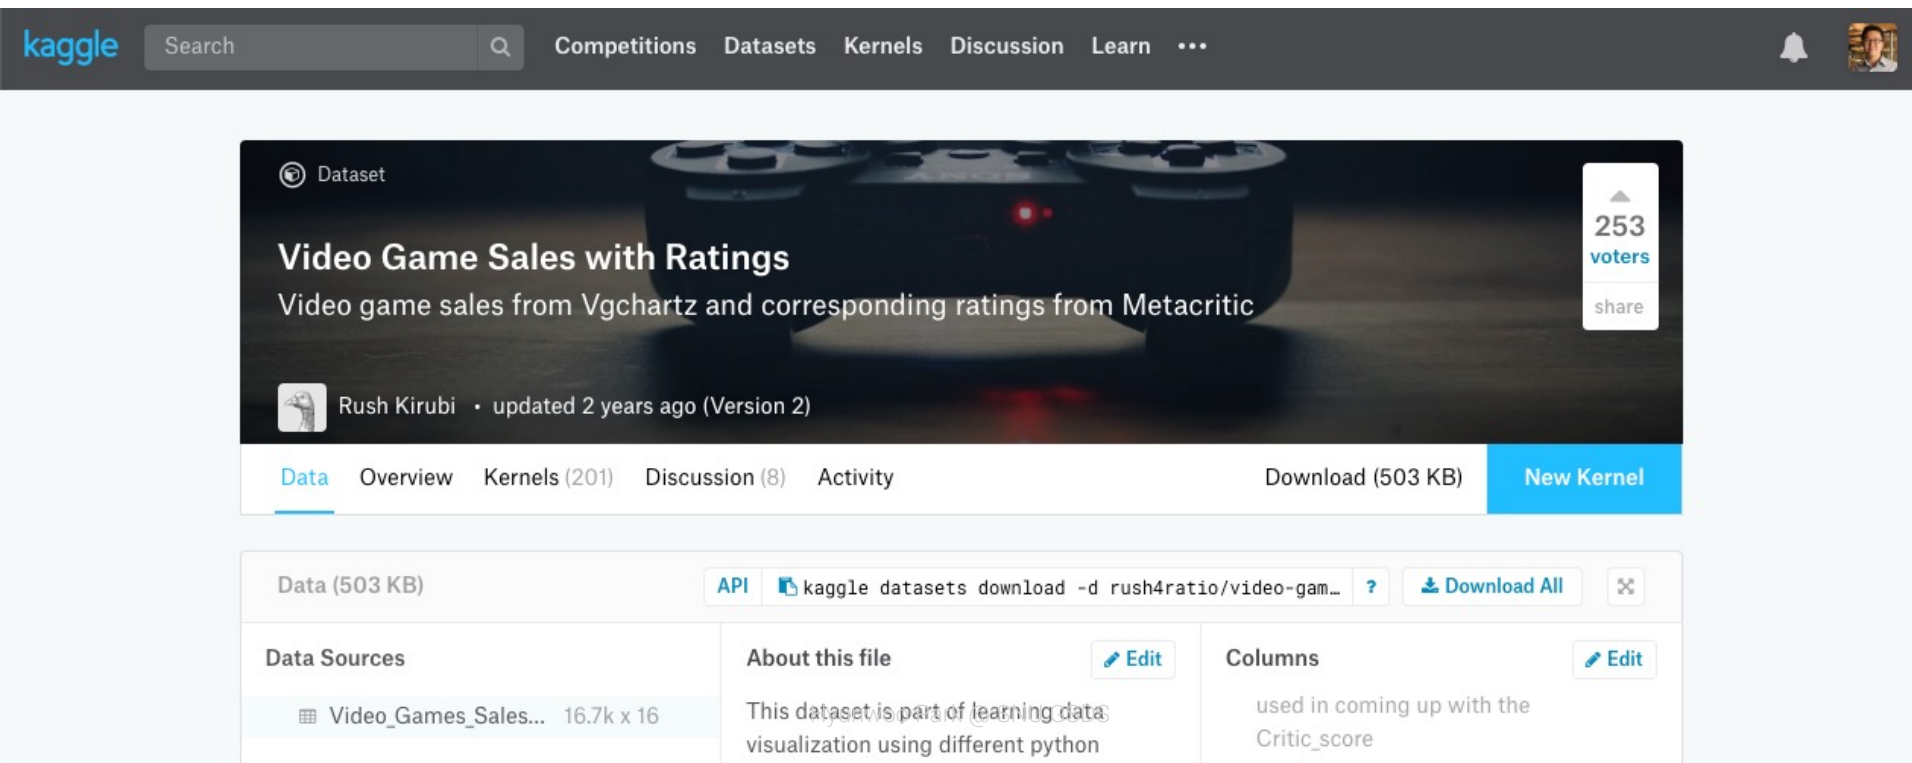
- Setup
 - Data Loading
 - Tidy Data
 - Data Transformation
- Descriptive Statistics
 - Describing One Variable
 - Types of Variables
 - Distributions
 - Summary Statistics
- Homework 1
 - Due 9/28 Tue Before Class 12:30pm
- Things to do
 - Review pandas / csv packages.
 - Compute summary statistics

Setup

Let's grab a dataset.

- Video Game Sales with Ratings

<https://www.kaggle.com/rush4ratio/video-game-sales-with-ratings>




The screenshot shows the Kaggle website interface. At the top is a dark navigation bar with the Kaggle logo, a search bar, and links for Competitions, Datasets, Kernels, Discussion, and Learn. Below this is the dataset page for 'Video Game Sales with Ratings' by Rush Kirubi. The page features a large header image of a video game controller with the dataset title and description. To the right of the header, it shows 253 voters and a share button. Below the header is a tabbed interface with 'Data' selected, and buttons for 'Download (503 KB)' and 'New Kernel'. The main content area is divided into three sections: 'Data Sources' showing the dataset file 'Video_Games_Sales...' with dimensions 16.7k x 16; 'About this file' with an 'Edit' button and a description stating it's part of a learning data visualization project; and 'Columns' with an 'Edit' button and a list of columns including 'Year', 'Genre', 'Platform', 'Sales', and 'Critic_score'.

Kaggle Search Competitions Datasets Kernels Discussion Learn ...

Dataset

Video Game Sales with Ratings

Video game sales from Vgchartz and corresponding ratings from Metacritic

 Rush Kirubi · updated 2 years ago (Version 2)

253 voters share

Data Overview Kernels (201) Discussion (8) Activity Download (503 KB) **New Kernel**

Data (503 KB) **API** `kaggle datasets download -d rush4ratio/video-gam...` **Download All**

Data Sources

Video_Games_Sales...	16.7k x 16
----------------------	------------

About this file

Edit

This dataset is part of learning data visualization using different python

Columns

Edit

used in coming up with the Critic_score

Data Loading

- Using Stata

1 insheet using Video_Games_Sales_as_at_22_Dec_2016.csv, c

The screenshot shows the Stata/MP 17.0 software interface. The top menu bar includes options like Log, Viewer, Graph, Do-file Editor, Data Editor, More, and Break. A search bar is located on the right. The main window is divided into two panes. The left pane, titled 'Results', displays the Stata logo, version information (17.0 MP-Parallel Edition), copyright details (1985-2021 StataCorp LLC), and the license information (Single-user 24-core perpetual). It also lists the serial number (501706317931) and the user (Hyunwoo Park, The Ohio State University). Below this, it shows the command entered: `. insheet using Video_Games_Sales_as_at_22_Dec_2016.csv, c` and the resulting output: (16 vars, 16,719 obs). The right pane, titled 'Variables', shows a list of variables with their names and labels. The variables listed are: name, platform, year_of_release, genre, publisher, na_sales, and eu_sales. The 'Properties' pane on the right shows the variable 'name' selected, with its properties (Name, Label, Type, Format, Value label, Notes) displayed.

Stata/MP 17.0

Log Viewer Graph Do-file Editor Data Editor More Break

Search

Results

Statistics and Data Science

Copyright 1985-2021 StataCorp LLC

StataCorp

4905 Lakeway Drive

College Station, Texas 77845 USA

800-STATA-PC <https://www.stata.com>

979-696-4600 stata@stata.com

Stata license: Single-user 24-core perpetual

Serial number: 501706317931

Licensed to: Hyunwoo Park

The Ohio State University

Notes:

1. Unicode is supported; see [help unicode_advice](#).
2. More than 2 billion observations are allowed; see [help obs_advice](#).
3. Maximum number of variables is set to 5,000; see [help set_maxvar](#).

. insheet using Video_Games_Sales_as_at_22_Dec_2016.csv, c

(16 vars, 16,719 obs)

Variables

Name	Label
name	Name
platform	Platform
year_of_release	Year_of_Release
genre	Genre
publisher	Publisher
na_sales	NA_Sales
eu_sales	EU_Sales

Properties

Variables

Name	Label
name	Name
platform	Platform
year_of_release	Year_of_Release
genre	Genre
publisher	Publisher
na_sales	NA_Sales
eu_sales	EU_Sales

Data

Frame	default
Filename	
Label	
Notes	
Variables	16
Observations	16,719

Data loading

- Using Stata

1 describe


The screenshot shows the Stata/MP 17.0 interface. The main window displays the results of the command `. describe`. The output indicates that the dataset contains 16,719 observations and 16 variables. Below this, a table lists the variables with their storage types, display formats, and labels.

Variable name	Storage type	Display format	Value label	Variable label
name	str132	%132s		Name
platform	str4	%9s		Platform
year_of_release	str4	%9s		Year_of_Release
genre	str12	%12s		Genre
publisher	str38	%38s		Publisher
na_sales	float	%9.0g		NA_Sales
eu_sales	float	%9.0g		EU_Sales
jp_sales	float	%9.0g		JP_Sales
other_sales	float	%9.0g		Other_Sales
global_sales	float	%9.0g		Global_Sales
critic_score	byte	%8.0g		Critic_Score
critic_count	int	%8.0g		Critic_Count
user_score	str3	%9s		User_Score
user_count	int	%8.0g		User_Count
developer	str80	%80s		Developer

The right-hand side of the interface features two panels. The 'Variables' panel lists the variables and their labels. The 'Properties' panel shows the current variable's properties, including its name, label, type, format, value label, and notes.

What is tidy data?

- It depends on the data context.
- Variables are easier to be linked; observations are harder.
- Further explanation on tidy data:
Wickham, H. (2014). Tidy data. Journal of Statistical Software, 59.
<http://vita.had.co.nz/papers/tidy-data.pdf>



More on tidy data

- Why are these not tidy?

religion	<\$10k	\$10-20k	\$20-30k	\$30-40k	\$40-50k	\$50-75k
Agnostic	27	34	60	81	76	137
Atheist	12	27	37	52	35	70
Buddhist	27	21	30	34	33	58
Catholic	418	617	732	670	638	1116
Don't know/refused	15	14	15	11	10	35
Evangelical Prot	575	869	1064	982	881	1486
Hindu	1	9	7	9	11	34
Historically Black Prot	228	244	236	238	197	223
Jehovah's Witness	20	27	24	24	21	30
Jewish	19	19	25	25	30	95

Table 4: The first ten rows of data on income and religion from the Pew Forum. Three columns, \$75-100k, \$100-150k and >150k, have been omitted

year	artist	track	time	date.entered	wk1	wk2	wk3
2000	2 Pac	Baby Don't Cry	4:22	2000-02-26	87	82	72
2000	2Ge+her	The Hardest Part Of ...	3:15	2000-09-02	91	87	92
2000	3 Doors Down	Kryptonite	3:53	2000-04-08	81	70	68
2000	98°0	Give Me Just One Nig...	3:24	2000-08-19	51	39	34
2000	A*Teens	Dancing Queen	3:44	2000-07-08	97	97	96
2000	Aaliyah	I Don't Wanna	4:15	2000-01-29	84	62	51
2000	Aaliyah	Try Again	4:03	2000-03-18	59	53	38
2000	Adams, Yolanda	Open My Heart	5:30	2000-08-26	76	76	74

Table 7: The first eight Billboard top hits for 2000. Other columns not shown are wk4, wk5, ..., wk75.

More on tidy data

- Are they now?

religion	income	freq
Agnostic	<\$10k	27
Agnostic	\$10-20k	34
Agnostic	\$20-30k	60
Agnostic	\$30-40k	81
Agnostic	\$40-50k	76
Agnostic	\$50-75k	137
Agnostic	\$75-100k	122
Agnostic	\$100-150k	109
Agnostic	>150k	84
Agnostic	Don't know/refused	96

Table 6: The first ten rows of the tidied Pew survey dataset on income and religion. The `column` has been renamed to `income`, and `value` to `freq`.

year	artist	time	track	date	week	rank
2000	2 Pac	4:22	Baby Don't Cry	2000-02-26	1	87
2000	2 Pac	4:22	Baby Don't Cry	2000-03-04	2	82
2000	2 Pac	4:22	Baby Don't Cry	2000-03-11	3	72
2000	2 Pac	4:22	Baby Don't Cry	2000-03-18	4	77
2000	2 Pac	4:22	Baby Don't Cry	2000-03-25	5	87
2000	2 Pac	4:22	Baby Don't Cry	2000-04-01	6	94
2000	2 Pac	4:22	Baby Don't Cry	2000-04-08	7	99
2000	2Ge+her	3:15	The Hardest Part Of ...	2000-09-02	1	91
2000	2Ge+her	3:15	The Hardest Part Of ...	2000-09-09	2	87
2000	2Ge+her	3:15	The Hardest Part Of ...	2000-09-16	3	92
2000	3 Doors Down	3:53	Kryptonite	2000-04-08	1	81
2000	3 Doors Down	3:53	Kryptonite	2000-04-15	2	70
2000	3 Doors Down	3:53	Kryptonite	2000-04-22	3	68
2000	3 Doors Down	3:53	Kryptonite	2000-04-29	4	67
2000	3 Doors Down	3:53	Kryptonite	2000-05-06	5	66

Table 8: First fifteen rows of the tidied billboard dataset. The `date` column does not appear in the original table, but can be computed from `date.entered` and `week`.

Data Transformation

Inspect the data.

- Number of observations?
- Number of variables?
- What types of variables does it have?

```
1 count
2 count if platform=="PS4"
3 describe
4 edit
```

5 verbs of data transformation

- Column operations: `select` (keep/drop), `mutate` (generate/replace)
- Row operations: `filter` (keep/drop), `arrange` (gsort)
- Summarize: `summarize` (summarize)

- You select and mutate “variables”,
filter and arrange “observations”.
- `mutate` = add or alter columns
`arrange` = sort

```
pd.DataFrame(csvdata[1:], columns=csvdata[0])
```

	Name	Platform	Year_of_Release	Genre	Publisher	NA_Sales	EU_Sales	JP_Sales	Other_Sales	Global_Sales
0	Wii Sports	Wii	2006	Sports	Nintendo	41.36	28.96	3.77	8.45	82.53
1	Super Mario Bros.	NES	1985	Platform	Nintendo	29.08	3.58	6.81	0.77	40.24
2	Mario Kart Wii	Wii	2008	Racing	Nintendo	15.68	12.76	3.79	3.29	35.52
3	Wii Sports Resort	Wii	2009	Sports	Nintendo	15.61	10.93	3.28	2.95	32.77
4	Pokemon Red/Pokemon Blue	GB	1996	Role-Playing	Nintendo	11.27	8.89	10.22	1	31.37
...
16714	Samurai Warriors: Sanada Maru	PS3	2016	Action	Tecmo Koei	0	0	0.01	0	0.01

Select

- You select “columns” or “variables” NOT “rows” or “observations”.
- Let’s select these 5 columns:
Name, Platform, Year_of_Release, Genre, Global_Sales.

1	preserve
2	keep name platform year_of_release genre global_sales
3	restore

Mutate

- You add a new column as a combination (or operation) of other columns.
You can also alter a current column by creating a new column with the same name.
- Let's create a new column called:
 $\text{Total_Sales} = \text{NA_Sales} + \text{EU_Sales} + \text{JP_Sales} + \text{Other_Sales}.$

```
1 preserve
2 g total_sales = na_sales + eu_sales + jp_sales + other_sales
3 restore
```

Filter

- You can keep a subset of observations by filtering out others.
- Let's collect videos games released on PS4 or Xbox One.

```
1 preserve
2 keep if platform=="PS4" | platform=="X0ne"
3 count
4 restore
```

🔑 Arrange

- You sort rows (or observations) with one or more criteria.
- Let's sort the data according to the following criteria in order:
 - (1) in descending order of **Year_of_Release**
 - (2) in ascending order of **Platform**
 - (3) in ascending order of **Name**

```
1 preserve
2 gsort -year_of_release platform name
3 edit if year_of_release != "N/A"
4 restore
```

Descriptive Statistics

📖 A brief look at descriptive statistics and visualization

📄 Video_Games_Sales_as_at_22_Dec_2016.csv (503.05 KB)

16 of 16 columns



	▲ Name ▼ Name of the game	▲ Platform ▼ Console on which the game is running	# Year_of_Release ▼ Year of the game released	▲ Genre ▼ Game's category	▲ Publisher ▼ Publisher	# NA_Sales Game sales in North America (units)
	11562 unique values	PS2 13% DS 13% Other (29) 74%	 1.98k 2.02k	Action 20% Sports 14% Other (10) 66%	Electronic Arts 8% Activision 6% Other (580) 86%	 0
1	Wii Sports	Wii	2006	Sports	Nintendo	
2	Super Mario Bros.	NES	1985	Platform	Nintendo	
3	Mario Kart Wii	Wii	2008	Racing	Nintendo	
4	Wii Sports Resort	Wii	2009	Sports	Nintendo	
5	Pokemon Red/Pokemon Blue	GB	1996	Role-Playing	Nintendo	
6	Tetris	GB	1989	Puzzle	Nintendo	
7	Mario Superstar	PS	2006	Platform	Nintendo	

Types of variables

- Categorical variables (or qualitative variable)
 - Platform: PS4, XOne, ...
 - Genre: Action, Sports, ...
- Numerical variables (or quantitative variable)
 - Discrete variables
 - Year_of_Release: 2006, 1985, 2008, 2009, ...
 - Continuous variables
 - Global_Sales: 82.53, 40.24, 35.52, 32.77, ...
- Understanding a single individual variable is to understand how it “varies” over different observations or measurements.
- Figuring out how observations are “distributed” along each of these variables is a good starting point.

🔑 Figuring out the distribution

- In essence, it is about counting observations that fall into a certain range.
- For categorical variable (and discrete variable sometimes), it's straightforward. And it's called tabulation.
- For continuous variable (and discrete variable sometimes), it requires another step before tabulation: binning.

Variable type	How to figure out the distribution?	Visualization
Categorical	Tabulation	Bar / Column Chart
Discrete	Both	Both
Continuous	Binning + Tabulation	Histogram, Density Plot, Boxplot

📌 Tabulation for categorical variable

- You just need to count the number of observations for each category.
- Let's count with **tab** function.
- Since categories don't have intrinsic ordering, it's usually useful to sort them by count in descending order.

1	<code>tab platform</code>
2	<code>tab platform, sort</code>

📌 Binning + tabulation for continuous variable

- For a continuous variable, you first need to discretize the variable into ranges.
- Three binning methods
 - makes groups of the given **width**
 - makes **n** groups with equal range
 - makes **n** groups with (approximately) equal numbers of observations

⌘ (Binning +) tabulation for discrete variable

- For a discrete variable, you can tabulate with or without binning.

Summary statistics for numerical variables

- Measure of location
 - Single representative numbers: mean, median
- Measure of spread
 - range, inter-quartile range, standard deviation
- Measure of rank
 - Five number summary
 - Minimum
 - First quartile
 - Median (= second quartile)
 - Third quartile
 - Maximum

Five-number summary

From Wikipedia, the free encyclopedia

The **five-number summary** is a set of **descriptive statistics** that provides information about a dataset. It consists of the five most important sample **percentiles**:

1. the **sample minimum** (*smallest observation*)
2. the **lower quartile** or *first quartile*
3. the **median** (the middle value)
4. the **upper quartile** or *third quartile*
5. the **sample maximum** (largest observation)

https://en.wikipedia.org/wiki/Five-number_summary

Summarize

- The fifth verb in data transformation is **summarize**.

```
1  replace year_of_release = "" if year_of_release=="N/A"  
2  deststring year_of_release, replace  
3  
4  su year_of_release  
5  su year_of_release, d
```

Export data

1	outsheet using test.csv, c replace
---	------------------------------------

📌 Some other useful Stata commands & concepts

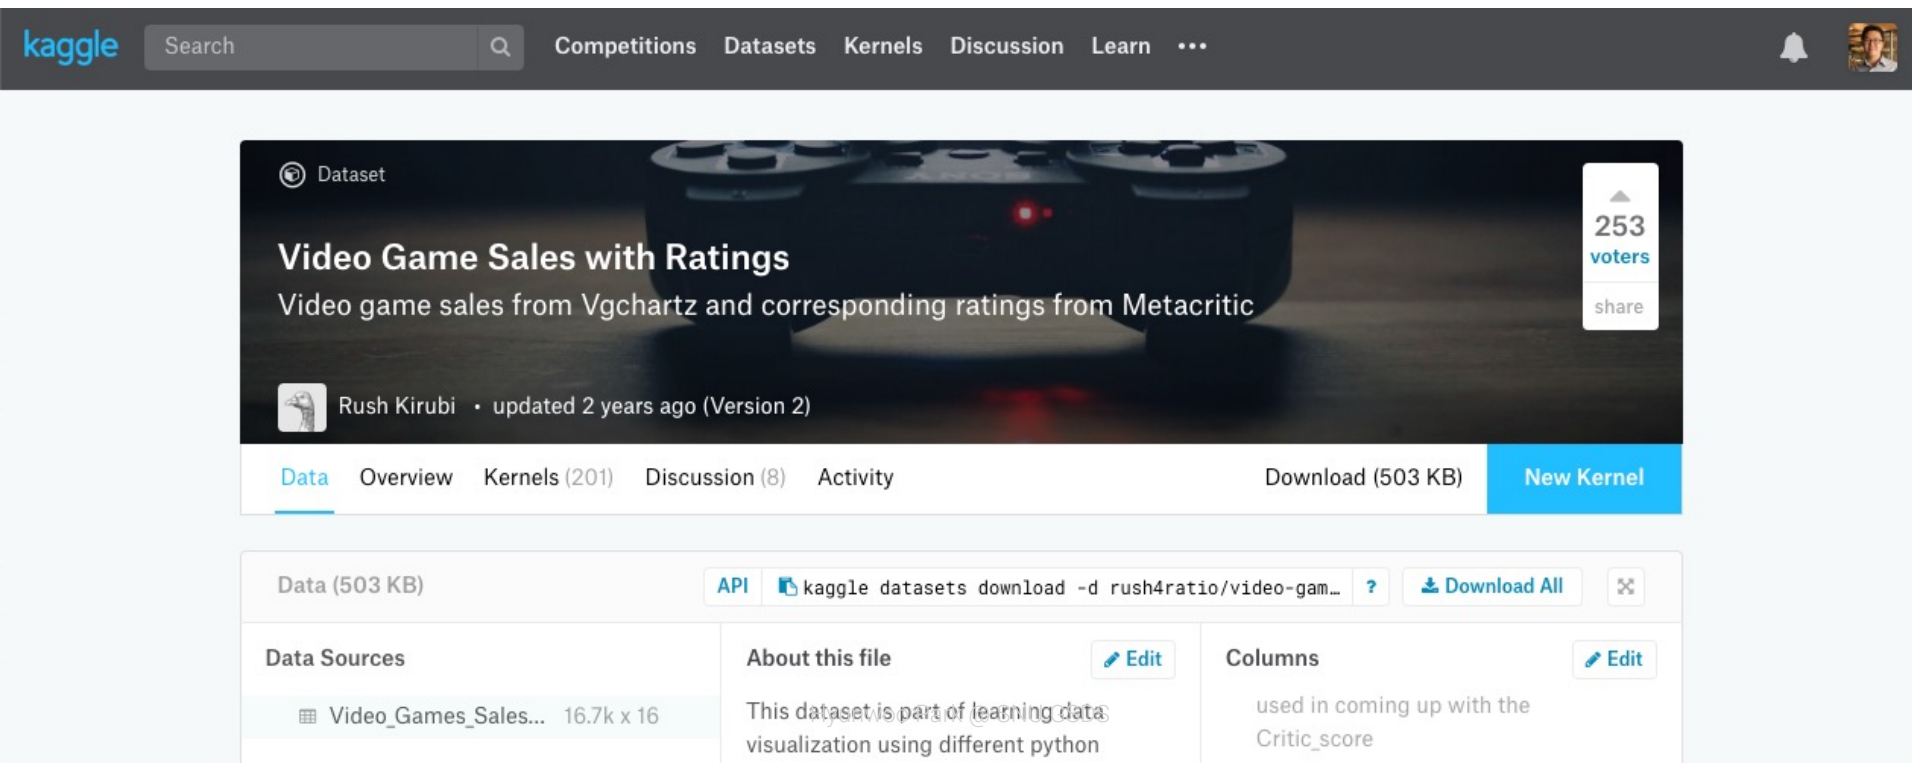
- egen (extensions to generate)
- collapse
- Do files
- Long form vs wide form

Appendix: Python-version

Let's grab a dataset.

- Video Game Sales with Ratings

<https://www.kaggle.com/rush4ratio/video-game-sales-with-ratings>




The screenshot shows the Kaggle website interface. At the top is a dark navigation bar with the Kaggle logo, a search bar, and links for Competitions, Datasets, Kernels, Discussion, and Learn. Below this is the dataset page for 'Video Game Sales with Ratings' by Rush Kirubi. The page features a large header image of a video game controller with the dataset title and description. On the right side of the header, it shows 253 voters and a share button. Below the header is a tabbed interface with 'Data' selected, and buttons for 'Download (503 KB)' and 'New Kernel'. The main content area is divided into three sections: 'Data Sources' showing the dataset file 'Video_Games_Sales...' with dimensions 16.7k x 16; 'About this file' with an 'Edit' button and a description stating it's part of a learning data visualization project; and 'Columns' with an 'Edit' button and a list of columns including 'Year', 'Genre', 'Platform', 'Sales', and 'Critic_score'.

Kaggle Search Competitions Datasets Kernels Discussion Learn ...

Dataset

Video Game Sales with Ratings

Video game sales from Vgchartz and corresponding ratings from Metacritic

 Rush Kirubi · updated 2 years ago (Version 2)

253 voters share

Data Overview Kernels (201) Discussion (8) Activity Download (503 KB) **New Kernel**

Data (503 KB) **API** `kaggle datasets download -d rush4ratio/video-gam...` **Download All**

Data Sources

Video_Games_Sales...	16.7k x 16
----------------------	------------

About this file

Edit

This dataset is part of learning data visualization using different python

Columns

Edit

used in coming up with the Critic_score

Data Loading

- Using pandas

```
import pandas as pd
pddata = pd.read_csv('Video_Games_Sales_as_at_22_Dec_2016.csv')
pddata
```

	Name	Platform	Year_of_Release	Genre	Publisher	NA_Sales	EU_Sales	JP_Sales	Other_Sales	Global_Sales	Critic_Score	Critic_Count	I
0	Wii Sports	Wii	2006.0	Sports	Nintendo	41.36	28.96	3.77	8.45	82.53	76.0	51.0	
1	Super Mario Bros.	NES	1985.0	Platform	Nintendo	29.08	3.58	6.81	0.77	40.24	NaN	NaN	
2	Mario Kart Wii	Wii	2008.0	Racing	Nintendo	15.68	12.76	3.79	3.29	35.52	82.0	73.0	
3	Wii Sports Resort	Wii	2009.0	Sports	Nintendo	15.61	10.93	3.28	2.95	32.77	80.0	73.0	
4	Pokemon Red/Pokemon Blue	GB	1996.0	Role-Playing	Nintendo	11.27	8.89	10.22	1.00	31.37	NaN	NaN	
...	
16714	Samurai Warriors: Sanada Maru	PS3	2016.0	Action	Tecmo Koei	0.00	0.00	0.01	0.00	0.01	NaN	NaN	
16715	LMA Manager 2007	X360	2006.0	Sports	Codemasters	0.00	0.01	0.00	0.00	0.01	NaN	NaN	

Data loading

- Using csv

```
import csv
csvdata = [r for r in csv.reader(open('Video_Games_Sales_as_at_22_Dec_2016.csv'))]
csvdata
```

```
[['Name',
  'Platform',
  'Year_of_Release',
  'Genre',
  'Publisher',
  'NA_Sales',
  'EU_Sales',
  'JP_Sales',
  'Other_Sales',
  'Global_Sales',
  'Critic_Score',
  'Critic_Count',
  'User_Score',
  'User_Count',
  'Developer',
  'Rating'],
 ['Wii Sports',
  'Wii',
  '2006',
  'Sports',
  'Nintendo',
  '41.26']]
```

What to use? pandas vs. csv

- pandas

- Numeric values are automatically parsed out.
- It returns a data frame object, which will work smoothly with numpy, jupyter notebook, and other numerical/scientific packages.
- Column-oriented data structure.
- [Pros] It comes with all the fancy helpers.
- [Pros] You benefit from performance improvements of the package.
- [Pros] You look like a data scientist.
- [Pros] Your data science friends/collaborators are likely using it.
- [Cons] It's not designed for row-wise operations.
- [Cons] Steeper learning curve

- CSV

- Everything is parsed as strings.
- It returns a list of lists.
- Row-oriented data structure.
- [Pros] Transparent
- [Pros] Easier to code for processing row by row
- [Cons] You have to do the housekeeping work yourself.
- [Cons] Easier to make mistake (You have to validate your data yourself.)
- [Cons] You look like a primate in data science.

🔗 Going back and forth between DataFrame and list of lists

- DataFrame to LoL

```
[list(pddata.columns)]+pddata.to_numpy().tolist()
```

```
[['Name',  
  'Platform',  
  'Year_of_Release',  
  'Genre',
```

- LoL to DataFrame

```
pd.DataFrame(csvdata[1:], columns=csvdata[0])
```


	Name	Platform	Year_of_Release	Genre	Publisher	NA_Sales	EU_Sales	JP_Sales	Other_Sales	Global_Sales
0	Wii Sports	Wii	2006	Sports	Nintendo	41.36	28.96	3.77	8.45	82.54
1	Super Mario Bros.	NES	1985	Platform	Nintendo	29.08	3.58	6.81	0.77	40.24
2	Mario Kart Wii	Wii	2008	Racing	Nintendo	15.68	12.76	3.79	3.29	35.52

📖 Your own library of helpers

- Add the folder containing this file to PYTHON_PATH.
(Or place it in the same folder of your working directory.)

```
1 import pandas as pd
2 import csv
3
4 def pd2csv(df): return [list(df.columns)]+df.to_numpy().tolist()
5 def csv2pd(lol): return pd.DataFrame(lol[1:], columns=lol[0])
6
7 def read_csv(fname):
8     print('Reading', fname)
9     return list(csv.reader(open(fname, 'r'))))
10
11 def write_csv(fname, data):
12     print('Writing', fname)
13     output = csv.writer(open(fname, 'w'))
14     output.writerows(data)
```

📖 Resources for pandas




- <https://github.com/wesm/pydata-book>

IPython Notebooks:

- Chapter 2: Python Language Basics, IPython, and Jupyter Notebooks
- Chapter 3: Built-in Data Structures, Functions, and Files
- Chapter 4: NumPy Basics: Arrays and Vectorized Computation
- Chapter 5: Getting Started with pandas
- Chapter 6: Data Loading, Storage, and File Formats
- Chapter 7: Data Cleaning and Preparation
- Chapter 8: Data Wrangling: Join, Combine, and Reshape
- Chapter 9: Plotting and Visualization
- Chapter 10: Data Aggregation and Group Operations
- Chapter 11: Time Series
- Chapter 12: Advanced pandas
- Chapter 13: Introduction to Modeling Libraries in Python
- Chapter 14: Data Analysis Examples
- Appendix A: Advanced NumPy

What is tidy data?

- It depends on the data context.
- Variables are easier to be linked; observations are harder.
- Further explanation on tidy data:
Wickham, H. (2014). Tidy data. Journal of Statistical Software, 59.
<http://vita.had.co.nz/papers/tidy-data.pdf>



More on tidy data

- Why are these not tidy?

religion	<\$10k	\$10-20k	\$20-30k	\$30-40k	\$40-50k	\$50-75k
Agnostic	27	34	60	81	76	137
Atheist	12	27	37	52	35	70
Buddhist	27	21	30	34	33	58
Catholic	418	617	732	670	638	1116
Don't know/refused	15	14	15	11	10	35
Evangelical Prot	575	869	1064	982	881	1486
Hindu	1	9	7	9	11	34
Historically Black Prot	228	244	236	238	197	223
Jehovah's Witness	20	27	24	24	21	30
Jewish	19	19	25	25	30	95

Table 4: The first ten rows of data on income and religion from the Pew Forum. Three columns, \$75-100k, \$100-150k and >150k, have been omitted

year	artist	track	time	date.entered	wk1	wk2	wk3
2000	2 Pac	Baby Don't Cry	4:22	2000-02-26	87	82	72
2000	2Ge+her	The Hardest Part Of ...	3:15	2000-09-02	91	87	92
2000	3 Doors Down	Kryptonite	3:53	2000-04-08	81	70	68
2000	98°0	Give Me Just One Nig...	3:24	2000-08-19	51	39	34
2000	A*Teens	Dancing Queen	3:44	2000-07-08	97	97	96
2000	Aaliyah	I Don't Wanna	4:15	2000-01-29	84	62	51
2000	Aaliyah	Try Again	4:03	2000-03-18	59	53	38
2000	Adams, Yolanda	Open My Heart	5:30	2000-08-26	76	76	74

Table 7: The first eight Billboard top hits for 2000. Other columns not shown are wk4, wk5, ..., wk75.

More on tidy data

- Are they now?

religion	income	freq
Agnostic	<\$10k	27
Agnostic	\$10-20k	34
Agnostic	\$20-30k	60
Agnostic	\$30-40k	81
Agnostic	\$40-50k	76
Agnostic	\$50-75k	137
Agnostic	\$75-100k	122
Agnostic	\$100-150k	109
Agnostic	>150k	84
Agnostic	Don't know/refused	96

Table 6: The first ten rows of the tidied Pew survey dataset on income and religion. The `column` has been renamed to `income`, and `value` to `freq`.

year	artist	time	track	date	week	rank
2000	2 Pac	4:22	Baby Don't Cry	2000-02-26	1	87
2000	2 Pac	4:22	Baby Don't Cry	2000-03-04	2	82
2000	2 Pac	4:22	Baby Don't Cry	2000-03-11	3	72
2000	2 Pac	4:22	Baby Don't Cry	2000-03-18	4	77
2000	2 Pac	4:22	Baby Don't Cry	2000-03-25	5	87
2000	2 Pac	4:22	Baby Don't Cry	2000-04-01	6	94
2000	2 Pac	4:22	Baby Don't Cry	2000-04-08	7	99
2000	2Ge+her	3:15	The Hardest Part Of ...	2000-09-02	1	91
2000	2Ge+her	3:15	The Hardest Part Of ...	2000-09-09	2	87
2000	2Ge+her	3:15	The Hardest Part Of ...	2000-09-16	3	92
2000	3 Doors Down	3:53	Kryptonite	2000-04-08	1	81
2000	3 Doors Down	3:53	Kryptonite	2000-04-15	2	70
2000	3 Doors Down	3:53	Kryptonite	2000-04-22	3	68
2000	3 Doors Down	3:53	Kryptonite	2000-04-29	4	67
2000	3 Doors Down	3:53	Kryptonite	2000-05-06	5	66

Table 8: First fifteen rows of the tidied billboard dataset. The `date` column does not appear in the original table, but can be computed from `date.entered` and `week`.

Data Transformation

Inspect the data.

- Number of observations?
- Number of variables?
- What types of variables does it have?

```
pd.DataFrame(csvdata[1:], columns=csvdata[0])
```

4	Red/Pokemon Blue	GB	1996	Role-Playing
...
16714	Samurai Warriors: Sanada Maru	PS3	2016	Action
16715	LMA Manager 2007	X360	2006	Sports
16716	Haitaka no Psychedelica	PSV	2016	Adventure
16717	Spirits & Spells	GBA	2003	Platform
16718	Winning Post 8 2016	PSV	2016	Simulation

16719 rows x 16 columns

```
for c in pddata.columns:  
    print([c])  
    print(pddata[c].describe(), '\n')
```

```
['Name']  
count          16717  
unique          11562  
top      Need for Speed: Most Wanted  
freq           12  
Name: Name, dtype: object
```

```
['Platform']  
count          16719  
unique           31  
top           PS2  
freq          2161  
Name: Platform, dtype: object
```

```
['Year_of_Release']  
count    16450.000000  
mean      2006.487356  
std         5.878995  
min      1980.000000  
25%      2003.000000  
50%      2007.000000  
75%      2010.000000  
max      2020.000000  
Name: Year_of_Release, dtype: float64
```


5 verbs of data transformation in tidyverse

- Column operations: `select`, `mutate`
- Row operations: `filter`, `arrange`
- Summarize: `summarize`

- You select and mutate “variables”,
filter and arrange “observations”.
- `mutate` = add or alter columns
`arrange` = sort

```
pd.DataFrame(csvdata[1:], columns=csvdata[0])
```

	Name	Platform	Year_of_Release	Genre	Publisher	NA_Sales	EU_Sales	JP_Sales	Other_Sales	Global_Sales
0	Wii Sports	Wii	2006	Sports	Nintendo	41.36	28.96	3.77	8.45	82.53
1	Super Mario Bros.	NES	1985	Platform	Nintendo	29.08	3.58	6.81	0.77	40.24
2	Mario Kart Wii	Wii	2008	Racing	Nintendo	15.68	12.76	3.79	3.29	35.52
3	Wii Sports Resort	Wii	2009	Sports	Nintendo	15.61	10.93	3.28	2.95	32.77
4	Pokemon Red/Pokemon Blue	GB	1996	Role-Playing	Nintendo	11.27	8.89	10.22	1	31.37
...
16714	Samurai Warriors: Sanada Maru	PS3	2016	Action	Tecmo Koei	0	0	0.01	0	0.01

Select

- You select “columns” or “variables” NOT “rows” or “observations”.
- Let’s select these 5 columns:

Name, Platform, Year_of_Release, Genre, Global_Sales.

```
pddata['Name, Platform, Year_of_Release, Genre, Global_Sales'].split(', ')]
```

	Name	Platform	Year_of_Release	Genre	Global_Sales
0	Wii Sports	Wii	2006.0	Sports	82.53
1	Super Mario Bros.	NES	1985.0	Platform	40.24
2	Mario Kart Wii	Wii	2008.0	Racing	35.52
3	Wii Sports Resort	Wii	2009.0	Sports	32.77
4	Pokemon Red/Pokemon Blue	GB	1996.0	Role-Playing	31.37
...
16714	Samurai Warriors: Sanada Maru	PS3	2016.0	Action	0.01
16715	LMA Manager 2007	X360	2006.0	Sports	0.01
16716	Haitaka no Psychedelica	PSV	2016.0	Adventure	0.01
16717	Spirits & Spells	GBA	2003.0	Platform	0.01
16718	Winning Post 8 2016	PSV	2016.0	Simulation	0.01

16719 rows x 5 columns

Mutate

- You add a new column as a combination (or operation) of other columns.
You can also alter a current column by creating a new column with the same name.
- Let's create a new column called:
$$\text{Total_Sales} = \text{NA_Sales} + \text{EU_Sales} + \text{JP_Sales} + \text{Other_Sales}.$$

```
pddata['Total_Sales'] = pddata['NA_Sales'] + pddata['EU_Sales'] + pddata['JP_Sales'] + pddata['Other_Sales']  
pddata
```

genre	Publisher	NA_Sales	EU_Sales	JP_Sales	Other_Sales	Global_Sales	Critic_Score	Critic_Count	User_Score	User_Count	Developer	Rating	Total_Sales
Sports	Nintendo	41.36	28.96	3.77	8.45	82.53	76.0	51.0	8	322.0	Nintendo	E	82.54
Form	Nintendo	29.08	3.58	6.81	0.77	40.24	NaN	NaN	NaN	NaN	NaN	NaN	40.24
Iceing	Nintendo	15.68	12.76	3.79	3.29	35.52	82.0	73.0	8.3	709.0	Nintendo	E	35.52
Sports	Nintendo	15.61	10.93	3.28	2.95	32.77	80.0	73.0	8	192.0	Nintendo	E	32.77
Role-playing	Nintendo	11.27	8.89	10.22	1.00	31.37	NaN	NaN	NaN	NaN	NaN	NaN	31.38
...
Action	Tecmo Koei	0.00	0.00	0.01	0.00	0.01	NaN	NaN	NaN	NaN	NaN	NaN	0.01

Filter

- You can keep a subset of observations by filtering out others.
- Let's collect videos games released on PS4 or Xbox One.

```
pddata[pddata.Platform.isin('PS4 XOne'.split())]
```

	Name	Platform	Year_of_Release	Genre	Publisher	NA_Sales	EU_Sales	JP_Sales	Other_Sales	Global_Sales	Critic_Score	Critic_Count	User_
31	Call of Duty: Black Ops 3	PS4	2015.0	Shooter	Activision	6.03	5.86	0.36	2.38	14.63	NaN	NaN	
42	Grand Theft Auto V	PS4	2014.0	Action	Take-Two Interactive	3.96	6.31	0.38	1.97	12.61	97.0	66.0	
77	FIFA 16	PS4	2015.0	Sports	Electronic Arts	1.12	6.12	0.06	1.28	8.57	82.0	42.0	
87	Star Wars Battlefront (2015)	PS4	2015.0	Shooter	Electronic Arts	2.99	3.49	0.22	1.28	7.98	NaN	NaN	
92	Call of Duty: Advanced Warfare	PS4	2014.0	Shooter	Activision	2.81	3.48	0.14	1.23	7.66	83.0	39.0	
...	
16634	Sébastien Loeb Rally Evo	XOne	2016.0	Racing	Milestone S.r.l	0.00	0.01	0.00	0.00	0.01	63.0	8.0	

Arrange

- You sort rows (or observations) with one or more criteria.
- Let's sort the data according to the following criteria in order:
 - (1) in descending order of **Year_of_Release**
 - (2) in ascending order of **Platform**
 - (3) in ascending order of **Name**

```
pd.data.sort_values('Year_of_Release Platform Name'.split(), ascending=[0,1,1])
```

	Name	Platform	Year_of_Release	Genre	Publisher	NA_Sales	EU_Sales	JP_Sales	Other_Sales	Global_Sales	Critic_Score	Critic_Count	Use
5936	Imagine: Makeup Artist	DS	2020.0	Simulation	Ubisoft	0.27	0.00	0.00	0.02	0.29	NaN	NaN	
14086	Phantasy Star Online 2 Episode 4: Deluxe Package	PS4	2017.0	Role-Playing	Sega	0.00	0.00	0.04	0.00	0.04	NaN	NaN	
16385	Brothers Conflict: Precious Baby	PSV	2017.0	Action	Idea Factory	0.00	0.00	0.01	0.00	0.01	NaN	NaN	
16222	Phantasy Star Online 2 Episode 4: Deluxe Package	PSV	2017.0	Role-Playing	Sega	0.00	0.00	0.01	0.00	0.01	NaN	NaN	
14985	Beyblade Burst	3DS	2016.0	Role-Playing	FuRyu	0.00	0.00	0.03	0.00	0.03	NaN	NaN	
...	
12785	Tom Clancy's Rainbow Six: Critical Hour	XB	NaN	Shooter	Unknown	0.04	0.01	0.00	0.00	0.06	54.0	10.0	

Same transformations with list of lists from csv

- Use just whatever way more intuitive and convenient for you.

```
1 # select
2 [[r[c] for c in [0,1,2,3,9]] for r in csvdata]
3
4 # mutate
5 for i, r in enumerate(csvdata):
6     if i==0: r.append('Total_Sales'); continue
7     r.append(sum([float(r[c]) for c in range(5,9)]))
8
9 # filter
10 [r for i, r in enumerate(csvdata) if i==0 or r[1] in 'PS4 XOne'.split()]
11
12 # arrange
13 keys = [(2,1), (1,0), (0,0)]
14 temp = csvdata
15 for k in reversed(keys):
16     temp = [temp[0]]+sorted(temp[1:], key=lambda r: r[k[0]], reverse=k[1])
```

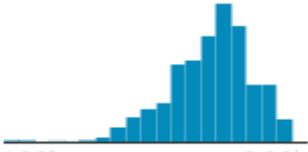

Descriptive Statistics

A brief look at descriptive statistics and visualization

Video_Games_Sales_as_at_22_Dec_2016.csv (503.05 KB)

16 of 16 columns



	<div>▲ Name ▼</div> <div>Name of the game</div>	<div>▲ Platform ▼</div> <div>Console on which the game is running</div>	<div># Year_of_Release ▼</div> <div>Year of the game released</div>	<div>▲ Genre ▼</div> <div>Game's category</div>	<div>▲ Publisher ▼</div> <div>Publisher</div>	<div># NA_Sales ▼</div> <div>Game sales in North America (units)</div>
	11562 unique values	<div>PS213%</div> <div>DS13%</div> <div>Other (29)74%</div>	<div>1.98k2.02k</div> 	<div>Action20%</div> <div>Sports14%</div> <div>Other (10)66%</div>	<div>Electronic Arts8%</div> <div>Activision6%</div> <div>Other (580)86%</div>	
1	Wii Sports	Wii	2006	Sports	Nintendo	
2	Super Mario Bros.	NES	1985	Platform	Nintendo	
3	Mario Kart Wii	Wii	2008	Racing	Nintendo	
4	Wii Sports Resort	Wii	2009	Sports	Nintendo	
5	Pokemon Red/Pokemon Blue	GB	1996	Role-Playing	Nintendo	
6	Tetris	GB	1989	Puzzle	Nintendo	
7	Mario Superstar	PS	2006	Platform	Nintendo	

Types of variables

- Categorical variables (or qualitative variable)
 - Platform: PS4, XOne, ...
 - Genre: Action, Sports, ...
- Numerical variables (or quantitative variable)
 - Discrete variables
 - Year_of_Release: 2006, 1985, 2008, 2009, ...
 - Continuous variables
 - Global_Sales: 82.53, 40.24, 35.52, 32.77, ...
- Understanding a single individual variable is to understand how it “varies” over different observations or measurements.
- Figuring out how observations are “distributed” along each of these variables is a good starting point.

🔑 Figuring out the distribution

- In essence, it is about counting observations that fall into a certain range.
- For categorical variable (and discrete variable sometimes), it's straightforward. And it's called tabulation.
- For continuous variable (and discrete variable sometimes), it requires another step before tabulation: binning.

Variable type	How to figure out the distribution?	Visualization
Categorical	Tabulation	Bar / Column Chart
Discrete	Both	Both
Continuous	Binning + Tabulation	Histogram, Density Plot, Boxplot

📌 Tabulation for categorical variable

- You just need to count the number of observations for each category.
- Let's count with `value_counts` function.
- Since categories don't have intrinsic ordering, it's usually useful to sort them by count in descending order.

```
pddata.Platform.value_counts()
```

PS2	2161
DS	2152
PS3	1331
Wii	1320
X360	1262
PSP	1209
PS	1197
PC	974
XB	824
GBA	822
GC	556
3DS	520
PSV	432
PS4	393
N64	319
XOne	247
SNES	239
SAT	173
WiiU	147

🔪 Binning + tabulation for continuous variable

- For a continuous variable, you first need to discretize the variable into ranges.
- Three binning methods
 - makes groups of the given `width`
 - makes `n` groups with equal range
 - makes `n` groups with (approximately) equal numbers of observations

```
1 pddata['Global_Sales'].value_counts(bins=10, sort=False)
2 pddata['Global_Sales'].value_counts(bins=pd.interval_range(start=-5, end=85,
3 freq=10), sort=False)
3 pd.qcut(pddata['Global_Sales'], q=10).value_counts(sort=False)
```

(-0.0735, 8.262]	16638	(-5, 5]	16512	(0.009000000000000001, 0.02]	1725
(8.262, 16.514]	58	(5, 15]	179	(0.02, 0.05]	2134
(16.514, 24.766]	13	(15, 25]	18	(0.05, 0.08]	1594
(24.766, 33.018]	7	(25, 35]	7	(0.08, 0.11]	1271
(33.018, 41.27]	2	(35, 45]	2	(0.11, 0.17]	1793
(41.27, 49.522]	0	(45, 55]	0	(0.17, 0.25]	1620
(49.522, 57.774]	0	(55, 65]	0	(0.25, 0.38]	1624
(57.774, 66.026]	0	(65, 75]	0	(0.38, 0.6]	1636
(66.026, 74.278]	0	(75, 85]	1	(0.6, 1.2]	1651
(74.278, 82.53]	1			(1.2, 82.53]	1671

Name: Global_Sales, dtype: int64 Name: Global_Sales, dtype: int64 Name: Global_Sales, dtype: int64

📌 (Binning +) tabulation for discrete variable

- For a discrete variable, you can tabulate with or without binning.

```
pddata['Year_of_Release'].value_counts()
```

```
2008.0    1427
2009.0    1426
2010.0    1255
2007.0    1197
2011.0    1136
2006.0    1006
2005.0     939
2002.0     829
2003.0     775
2004.0     762
2012.0     653
2015.0     606
2014.0     581
2013.0     544
2016.0     502
2001.0     482
1998.0     379
```

```
pddata['Year_of_Release'].value_counts(
    bins=pd.interval_range(start=1975, end=2025, freq=10), sort=False)

(1975, 1985]    136
(1985, 1995]    571
(1995, 2005]   5406
(2005, 2015]   9831
(2015, 2025]    506
Name: Year_of_Release, dtype: int64
```

Summary statistics for numerical variables

- Measure of location
 - Single representative numbers: mean, median
- Measure of spread
 - range, inter-quartile range, standard deviation
- Measure of rank
 - Five number summary
 - Minimum
 - First quartile
 - Median (= second quartile)
 - Third quartile
 - Maximum

Five-number summary

From Wikipedia, the free encyclopedia

The **five-number summary** is a set of **descriptive statistics** that provides information about a dataset. It consists of the five most important sample **percentiles**:

1. the **sample minimum** (*smallest observation*)
2. the **lower quartile** or *first quartile*
3. the **median** (the middle value)
4. the **upper quartile** or *third quartile*
5. the **sample maximum** (largest observation)

https://en.wikipedia.org/wiki/Five-number_summary

📌 Summarize

- The fifth verb in data transformation is `summarize`,
In pandas, you use `describe`.

```
pddata[ 'Year_of_Release' ].describe()
```

```
count      16450.000000
mean       2006.487356
std         5.878995
min        1980.000000
25%        2003.000000
50%        2007.000000
75%        2010.000000
max         2020.000000
Name: Year_of_Release, dtype: float64
```

Export data



[Getting started](#) [User Guide](#) [API reference](#) [Development](#) [Release notes](#)



🔍 Search the docs ...

Input/output

[pandas.read_pickle](#)
[pandas.DataFrame.to_pickle](#)
[pandas.read_table](#)
[pandas.read_csv](#)
[pandas.DataFrame.to_csv](#)
[pandas.read_fwf](#)
[pandas.read_clipboard](#)
[pandas.DataFrame.to_clipboard](#)
[pandas.read_excel](#)
[pandas.DataFrame.to_excel](#)
[pandas.ExcelFile.parse](#)
[pandas.io.formats.style.Styler.to_excel](#)
[pandas.ExcelWriter](#)
[pandas.io.json.read_json](#)
[pandas.io.json.to_json](#)
[pandas.io.json.build_table_schema](#)
[pandas.read_html](#)

Input/output

Pickling

`read_pickle(filepath_or_buffer[, ...])` Load pickled pandas object (or any object) from file.

`DataFrame.to_pickle(path[, compression, ...])` Pickle (serialize) object to file.

Flat file

`read_table(filepath_or_buffer[, sep, ...])` Read general delimited file into DataFrame.

`read_csv(filepath_or_buffer[, sep, ...])` Read a comma-separated values (csv) file into DataFrame.

`DataFrame.to_csv([path_or_buf, sep, na_rep, ...])` Write object to a comma-separated values (csv) file.

`read_fwf(filepath_or_buffer[, colspecs, ...])` Read a table of fixed-width formatted lines into DataFrame.

☰ On this page

[Pickling](#)
[Flat file](#)
[Clipboard](#)
[Excel](#)
[JSON](#)
[HTML](#)
[XML](#)
[Latex](#)
[HDFStore: PyTables \(HDF5\)](#)
[Feather](#)
[Parquet](#)
[ORC](#)
[SAS](#)
[SPSS](#)
[SQL](#)
[Google BigQuery](#)
[STATA](#)

Common discrete distributions

Distribution	R function	Example
Bernoulli	<code>rbernoulli</code>	Head or tail in a coin toss
Rectangular (or uniform)		Number of 1's in n dice rolls
Binomial	<code>rbinom</code>	Number of heads in n coin tosses
Geometric	<code>rgeom</code>	Number of failures until the first success
Negative binomial	<code>rnbinom</code>	Number of failures until nth success
Poisson	<code>rpois</code>	Number of events if occurrences are independent from each other
Zipf		Number of occurrences of words in texts

Common continuous distributions

Distribution	R function	Example
Uniform	<code>runif</code>	Random number within a range
Normal	<code>rnorm</code>	Height of people in a population
Exponential	<code>rexp</code>	Length of time between independent events
Lognormal	<code>rlnorm</code>	Length of comments, system repair times, income distribution
Pareto		Wealth distribution, city size, size of meteorites, 80/20 rule

📌 Create a new DataFrame with random variables

```
1  from numpy.random import default_rng
2  rng = default_rng()
3  N = 1000
4  data = [
5      rng.binomial(1,.1,size=N),
6      rng.integers(1,7,size=N),
7      rng.binomial(10,.1,size=N),
8      rng.geometric(.1,size=N),
9      rng.negative_binomial(10,.1,size=N),
10     rng.poisson(10,size=N),
11     rng.uniform(10,20,size=N),
12     rng.normal(10,2,size=N),
13     rng.exponential(.1,size=N),
14     rng.lognormal(2,2,size=N)
15 ]
16 cnames = 'Bernoulli RandInt Binomial Geometric NBinom Poisson Uniform Normal
17 Exp LogNormal'.split()
18 randdf = pd.DataFrame({cnames[i]: data[i] for i in range(len(data))})
randdf
```

- Try to compute the summary statistics table about this randomly generated dataset.